

NOLAN POZZOBON

EDUCATION

Carnegie Mellon University , Pittsburgh, PA <i>M.S. Artificial Intelligence and Innovation (MSAI)</i>	Expected May 2028
The University of Chicago , Chicago, IL <i>B.S. Computer Science (Machine Learning specialization); B.A. Linguistics — GPA: 3.81 / 4.00</i>	June 2026
Relevant Coursework: Algorithms, Discrete Mathematics, Distributed Systems, Database Systems, Machine Learning, Natural Language Processing, Math for Machine Learning, Large Language Models, Adversarial AI, Computational Linguistics Awards: Scholar All-American, NCAA Wrestling (2024, 2026)	

TECHNICAL SKILLS

Languages: Python, Rust, SQL, JavaScript
ML / AI: PyTorch, Hugging Face Transformers, vLLM, LoRA/PEFT, RAG, LangGraph/LangChain, LLM evaluation
Tools & Infrastructure: Git, Docker, Linux, AWS Lambda, Modal, Lambda Labs, PostgreSQL, Neo4j, Vue, NumPy, scikit-learn

EXPERIENCE

Opteryx — AI Engineering Fellow <i>Valencia, Spain (Hybrid)</i>	June 2026 – August 2026
- Shipped a document-humanization feature into a production LLM pipeline generating regulated pharmaceutical stability reports, delivering it to paying enterprise customers. - Built feature extraction over heterogeneous document corpora to model stylistic and structural patterns, and rearchitected the generation pipeline to consume it.	
UPS — Supply Chain Solutions Intern <i>Alpharetta, GA</i>	July 2025 – September 2025
- Applied mixed-integer linear programming and clustering algorithms to franchise region assignment across the US, reducing operating cost by \$100,000 annually. - Implemented a constrained KNN algorithm to group packages by ZIP code and size range, improving last-mile delivery routing efficiency nationwide.	
Toyota Technological Institute at Chicago — Research Assistant <i>Chicago, IL — Advisor: Dr. Zhewei Sun</i>	January 2025 – June 2025
- Extended and modified LLM architectures in PyTorch across distributed GPUs on the TTIC cluster, generalizing slang-embedding research to a multilingual setting. - Built web-scraping and cleaning pipelines to assemble novel multilingual datasets for fine-tuning.	
CharacTour — Software Engineering Intern	June 2024 – August 2024
Built generative-AI tooling to rewrite platform content for 300 distinct personality profiles, and refined the production recommendation matching algorithm.	
ESGX — Product Engineering Intern	August 2024 – January 2025

PROJECTS

Nabu — AI Language Tutor (Python, LangGraph, Whisper, ElevenLabs)	July 2025 – September 2025
Built an AI-native desktop application for real-time spoken conversation, designing the agent workflows, prompt orchestration, and API endpoints across three external speech and language services.	
Multi-Agent Claim Evaluation System (Python, RAG, NLI)	September 2025 – December 2025
- Engineered a multi-agent pipeline combining natural language inference, retrieval over embedded user histories, and external grounding to flag unsupported claims. - Benchmarked across Gemma 3, Llama 3.1, and GPT-OSS with comparative ablations and a human-annotated evaluation set, isolating retrieval quality as the primary bottleneck.	
Targeted Knowledge Updates in LLMs (PyTorch, LoRA)	September 2025 – December 2025
- Implemented inference-time control over the model-editing tradeoff via an alpha sweep on LoRA task vectors, achieving 100% edit efficacy at 87% broad task accuracy with no retraining. - Validated capability isolation with a 50-item multi-domain locality set, preserving 100% code accuracy and near-baseline perplexity.	

RESEARCH

“LLM Commitment Failure: Localizing Multilingual Sycophancy to Commitment Under Pressure” — sole author, under review at BlackboxNLP. Evaluated 11 models across 7 languages; localized the mechanism with linear probes, the logit lens, and activation patching.	
---	--